

```

!mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/

!kaggle datasets download -d samuelcortinhas/cats-and-dogs-image-classification

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /root/.kaggle/kaggle.js
Downloading cats-and-dogs-image-classification.zip to /content
 96% 62.0M/64.4M [00:02<00:00, 30.9MB/s]
100% 64.4M/64.4M [00:02<00:00, 23.8MB/s]

```

---

```

import zipfile
zip_ref = zipfile.ZipFile('/content/cats-and-dogs-image-classification.zip', 'r')
zip_ref.extractall('/content')
zip_ref.close()

import tensorflow as tf
from tensorflow import keras
from keras import Sequential
from keras.layers import Dense, Conv2D, MaxPooling2D, Flatten, BatchNormalization, Dropout

# generators
train_ds = keras.utils.image_dataset_from_directory(
    directory = '/content/train',
    labels='inferred',
    label_mode = 'int',
    batch_size=32,
    image_size=(256,256)
)

validation_ds = keras.utils.image_dataset_from_directory(
    directory = '/content/train',
    labels='inferred',
    label_mode = 'int',
    batch_size=32,
    image_size=(256,256)
)

    Found 557 files belonging to 2 classes.
    Found 557 files belonging to 2 classes.

# Normalize
def process(image,label):
    image = tf.cast(image/255. ,tf.float32)
    return image,label

train_ds = train_ds.map(process)
validation_ds = validation_ds.map(process)

# create CNN model

model = Sequential()

model.add(Conv2D(32, kernel_size=(3,3), padding='valid', activation='relu', input_shape=(256, 256, 3)))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))

model.add(Conv2D(64, kernel_size=(3,3), padding='valid', activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))

model.add(Conv2D(128, kernel_size=(3,3), padding='valid', activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))

model.add(Flatten())

model.add(Dense(128, activation='relu'))
model.add(Dropout(0.1))
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.1))
model.add(Dense(1, activation='sigmoid'))

model.summary()

Model: "sequential"

```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 254, 254, 32)	896
batch_normalization (Batch Normalization)	(None, 254, 254, 32)	128
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_1 (Conv2D)	(None, 125, 125, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 125, 125, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_2 (Conv2D)	(None, 60, 60, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 60, 60, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 128)	0
flatten (Flatten)	(None, 115200)	0
dense (Dense)	(None, 128)	14745728
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 64)	8256
dropout_1 (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 1)	65

=====  
 Total params: 14848193 (56.64 MB)  
 Trainable params: 14847745 (56.64 MB)  
 Non-trainable params: 448 (1.75 KB)

```
model.compile(optimizer='adam',loss='binary_crossentropy',metrics=['accuracy'])
```

```
history = model.fit(train_ds,epochs=10,validation_data=validation_ds)
```

```

Epoch 1/10
18/18 [=====] - 20s 265ms/step - loss: 6.3124 - accuracy: 0.5745 - val_loss: 3.4313 - val_accuracy: 0.4740
Epoch 2/10
18/18 [=====] - 7s 285ms/step - loss: 5.1244 - accuracy: 0.6445 - val_loss: 10.0341 - val_accuracy: 0.4991
Epoch 3/10
18/18 [=====] - 6s 296ms/step - loss: 2.2824 - accuracy: 0.7038 - val_loss: 12.8366 - val_accuracy: 0.4991
Epoch 4/10
18/18 [=====] - 6s 251ms/step - loss: 2.3799 - accuracy: 0.7038 - val_loss: 4.7074 - val_accuracy: 0.5045
Epoch 5/10
18/18 [=====] - 6s 261ms/step - loss: 1.8279 - accuracy: 0.7469 - val_loss: 14.9632 - val_accuracy: 0.5009
Epoch 6/10
18/18 [=====] - 5s 250ms/step - loss: 1.9379 - accuracy: 0.7846 - val_loss: 8.7747 - val_accuracy: 0.5117
Epoch 7/10
18/18 [=====] - 8s 429ms/step - loss: 1.3015 - accuracy: 0.7882 - val_loss: 6.2203 - val_accuracy: 0.5566
Epoch 8/10
18/18 [=====] - 5s 243ms/step - loss: 1.3880 - accuracy: 0.8097 - val_loss: 3.1116 - val_accuracy: 0.5943
Epoch 9/10
18/18 [=====] - 6s 286ms/step - loss: 1.2661 - accuracy: 0.8241 - val_loss: 1.3596 - val_accuracy: 0.6625
Epoch 10/10
18/18 [=====] - 5s 242ms/step - loss: 1.2965 - accuracy: 0.8187 - val_loss: 2.7764 - val_accuracy: 0.4937


```

```
import matplotlib.pyplot as plt
```


```

plt.plot(history.history['accuracy'],color='red',label='train')
plt.plot(history.history['val_accuracy'],color='blue',label='validation')
plt.legend()
plt.show()


```




```
plt.plot(history.history['accuracy'],color='red',label='train')
plt.plot(history.history['val_accuracy'],color='blue',label='validation')
plt.legend()
plt.show()
```



```
plt.plot(history.history['loss'],color='red',label='train')
plt.plot(history.history['val_loss'],color='blue',label='validation')
plt.legend()
plt.show()
```



```
plt.plot(history.history['loss'],color='red',label='train')
plt.plot(history.history['val_loss'],color='blue',label='validation')
plt.legend()
plt.show()
```



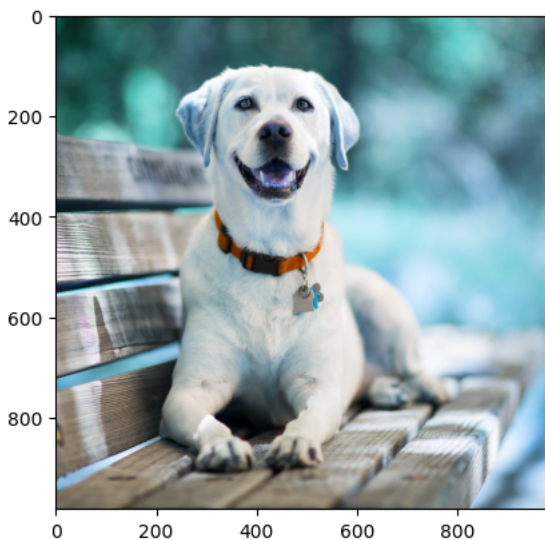
```
import cv2
2 1
test_img = cv2.imread('/content/cat.jpg')
plt.imshow(test_img)
test_img = cv2.resize(test_img, (256, 256))
test_input = test_img.reshape((1, 256, 256, 3)) / 255. # Normalize the test image
```



```
prediction = model.predict(test_input)
print("Prediction:", prediction)

1/1 [=====] - 1s 625ms/step
Prediction: [[0.20155276]]
```

```
test_img = cv2.imread('/content/dog.jpg')
plt.imshow(test_img)
test_img = cv2.resize(test_img, (256, 256))
test_input = test_img.reshape((1, 256, 256, 3)) / 255. # Normalize the test image
```



```
prediction = model.predict(test_input)
print("Prediction:", prediction)
```

```
1/1 [=====] - 0s 27ms/step
Prediction: [[0.08369035]]
```

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
history = model.fit(train_ds, epochs=20, validation_data=validation_ds)
```

```
Epoch 1/20
18/18 [=====] - 10s 257ms/step - loss: 1.5615 - accuracy: 0.8276 - val_loss: 19.2181 - val_accuracy: 0.5000
Epoch 2/20
18/18 [=====] - 7s 326ms/step - loss: 1.4516 - accuracy: 0.8151 - val_loss: 19.1029 - val_accuracy: 0.5009
Epoch 3/20
18/18 [=====] - 6s 264ms/step - loss: 1.2592 - accuracy: 0.8384 - val_loss: 23.4758 - val_accuracy: 0.5009
Epoch 4/20
18/18 [=====] - 5s 262ms/step - loss: 0.9870 - accuracy: 0.8743 - val_loss: 4.8519 - val_accuracy: 0.4991
Epoch 5/20
18/18 [=====] - 6s 264ms/step - loss: 0.8796 - accuracy: 0.8815 - val_loss: 16.9837 - val_accuracy: 0.4991
Epoch 6/20
18/18 [=====] - 5s 244ms/step - loss: 0.4551 - accuracy: 0.9246 - val_loss: 7.7307 - val_accuracy: 0.4991
Epoch 7/20
18/18 [=====] - 6s 264ms/step - loss: 0.3452 - accuracy: 0.9354 - val_loss: 9.1421 - val_accuracy: 0.4991
Epoch 8/20
18/18 [=====] - 6s 264ms/step - loss: 0.2563 - accuracy: 0.9479 - val_loss: 14.6302 - val_accuracy: 0.4991
Epoch 9/20
18/18 [=====] - 6s 273ms/step - loss: 0.3344 - accuracy: 0.9551 - val_loss: 2.0666 - val_accuracy: 0.5009
Epoch 10/20
18/18 [=====] - 6s 314ms/step - loss: 0.3430 - accuracy: 0.9605 - val_loss: 7.7986 - val_accuracy: 0.4991
Epoch 11/20
18/18 [=====] - 6s 264ms/step - loss: 0.1465 - accuracy: 0.9803 - val_loss: 19.6123 - val_accuracy: 0.5009
Epoch 12/20
18/18 [=====] - 6s 264ms/step - loss: 0.1410 - accuracy: 0.9767 - val_loss: 16.3161 - val_accuracy: 0.5009
Epoch 13/20
18/18 [=====] - 6s 319ms/step - loss: 0.1334 - accuracy: 0.9767 - val_loss: 19.5058 - val_accuracy: 0.5009
Epoch 14/20
18/18 [=====] - 6s 265ms/step - loss: 0.0733 - accuracy: 0.9838 - val_loss: 32.4718 - val_accuracy: 0.5009
Epoch 15/20
18/18 [=====] - 8s 413ms/step - loss: 0.2472 - accuracy: 0.9785 - val_loss: 25.5246 - val_accuracy: 0.5009
Epoch 16/20
18/18 [=====] - 6s 266ms/step - loss: 0.3528 - accuracy: 0.9497 - val_loss: 19.0979 - val_accuracy: 0.5009
Epoch 17/20
18/18 [=====] - 7s 322ms/step - loss: 0.2538 - accuracy: 0.9569 - val_loss: 33.0052 - val_accuracy: 0.5009
Epoch 18/20
18/18 [=====] - 5s 247ms/step - loss: 0.1665 - accuracy: 0.9695 - val_loss: 28.6201 - val_accuracy: 0.5009
Epoch 19/20
18/18 [=====] - 6s 264ms/step - loss: 0.2495 - accuracy: 0.9641 - val_loss: 31.7087 - val_accuracy: 0.5009
Epoch 20/20
18/18 [=====] - 7s 291ms/step - loss: 0.2037 - accuracy: 0.9767 - val_loss: 31.0354 - val_accuracy: 0.5009
```

```
# Class mapping
```

```
class_mapping = {0: 'Cat', 1: 'Dog'}
```

```
# Load test images
```

```
test_img_cat = cv2.imread('/content/cat.jpg')
```

```
test_img_cat = cv2.resize(test_img_cat, (256, 256))
```

```
test_input_cat = test_img_cat.reshape((1, 256, 256, 3)) / 255.
```

```
test_img_dog = cv2.imread('/content/dog.jpg')
```

```
test_img_dog = cv2.resize(test_img_dog, (256, 256))
```

```
test_input_dog = test_img_dog.reshape((1, 256, 256, 3)) / 255.
```

```
# Model prediction on test images
```

```
prediction_cat = model.predict(test_input_cat)
```

```
prediction_dog = model.predict(test_input_dog)
```

```
1/1 [=====] - 0s 359ms/step
1/1 [=====] - 0s 125ms/step
```

```
# Interpret predictions using class mapping
```

```
class_prediction_cat = class_mapping[int(round(prediction_cat[0][0]))]
```

```
class_prediction_dog = class_mapping[int(round(prediction_dog[0][0]))]
```

```
print("Prediction for cat:", class_prediction_cat)
```

```
print("Prediction for dog:", class_prediction_dog)
```

```
Prediction for cat: Cat
Prediction for dog: Cat

# Evaluate the model on a test set (assuming you have a separate test directory)
test_ds = keras.utils.image_dataset_from_directory(
    directory='/content/test',
    labels='inferred',
    label_mode='int',
    batch_size=32,
    image_size=(256, 256)
)

test_ds = test_ds.map(process)

test_loss, test_accuracy = model.evaluate(test_ds)
print(f'Test Accuracy: {test_accuracy}')
print(f'Test Loss: {test_loss}')

Found 140 files belonging to 2 classes.
5/5 [=====] - 1s 92ms/step - loss: 31.0931 - accuracy: 0.5000
Test Accuracy: 0.5
Test Loss: 31.093067169189453
```